

Safety & Red Teaming Annotation

Testing AI System Boundaries

Jin Zhao

Brandeis University

April 22, 2025

Today's Agenda

- ① Red teaming for AI systems
- ② Adversarial prompt annotation
- ③ Harmfulness and toxicity annotation
- ④ Content moderation datasets
- ⑤ Safety evaluation benchmarks
- ⑥ Ethical considerations in safety annotation

Note: This lecture discusses sensitive content categories

What is Red Teaming?

Adversarial testing of AI systems

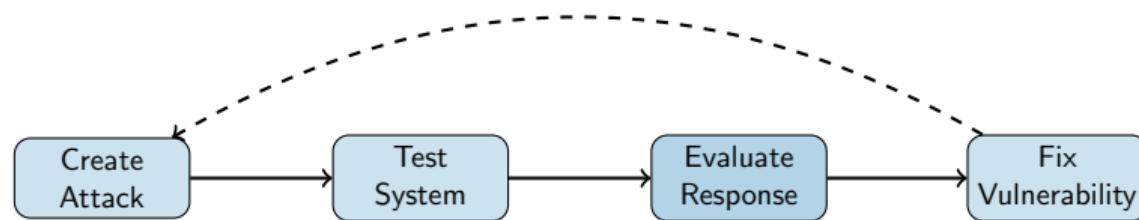
Origin: Military/security term for adversarial testing

For LLMs:

- Finding inputs that cause harmful outputs
- Testing system boundaries
- Identifying failure modes
- Improving safety before deployment

Goal: Break the system (constructively)

Red Teaming Workflow



Annotation role: Evaluating whether responses are harmful

Types of Adversarial Prompts

Categories red teamers explore:

① Jailbreaks: Bypass safety filters

- Role-playing scenarios
- Hypothetical framing
- Encoding tricks

② Harmful requests:

- Dangerous information
- Illegal activities
- Self-harm content

③ Bias elicitation:

- Stereotypes
- Discriminatory outputs

Two annotation tasks:

1. Attack classification:

- What type of attack is this prompt?
- What harm category does it target?
- How sophisticated is the attack?

2. Response evaluation:

- Did the model refuse appropriately?
- Was harmful content generated?
- Was the refusal too aggressive?

Harmfulness Taxonomy

Common harm categories:

Content harms:

- Violence
- Hate speech
- Sexual content
- Self-harm
- Misinformation

Operational harms:

- Privacy violations
- Security exploits
- Illegal advice
- Professional harm
- Manipulation

Challenge: Categories overlap and context matters

Severity Levels

Not all harms are equal

Common scale:

- ① **Safe:** No harmful content
- ② **Mild:** Minor issues, low risk
- ③ **Moderate:** Concerning content, medium risk
- ④ **Severe:** Dangerous content, high risk
- ⑤ **Critical:** Immediate danger, must block

Annotation challenge:

- Calibrating severity across annotators
- Context-dependent severity
- Cultural variation in perceived severity

Detecting toxic language

Definition: Rude, disrespectful, or unreasonable language

Sub-categories:

- **Insults:** Personal attacks
- **Threats:** Expressions of intent to harm
- **Profanity:** Obscene language
- **Identity attacks:** Targeting protected groups
- **Sexual harassment:** Unwanted sexual content

Tools: Perspective API, custom classifiers

Key datasets for safety research:

- **Civil Comments:** Toxicity labels on 2M comments
- **Jigsaw:** Multi-label toxicity annotations
- **HateXplain:** Hate speech with rationales
- **ToxiGen:** Machine-generated toxic statements
- **RealToxicityPrompts:** Prompt completion toxicity

Common features:

- Multiple annotators per item
- Demographic information
- Identity-based subcategories

Standardized safety evaluation:

- **TruthfulQA:** Testing for truthfulness
- **BBQ:** Bias benchmark with questions
- **WinoBias:** Gender bias in coreference
- **CrowS-Pairs:** Stereotype measurement
- **HarmBench:** Comprehensive harm evaluation

Annotation role:

- Creating gold standard responses
- Validating automated metrics
- Identifying new harm categories

Who should annotate safety content?

Considerations:

- **Training:** Extensive guidelines and calibration
- **Diversity:** Multiple perspectives on harm
- **Expertise:** Domain experts for specific harms
- **Wellbeing:** Mental health support required

Important:

- Exposure to harmful content is psychologically taxing
- Rotation and breaks are essential
- Counseling should be available

Ethical obligation to annotators

Risks:

- Secondary trauma from harmful content
- Desensitization over time
- Moral injury

Mitigations:

- Content warnings before annotation
- Session time limits
- Regular breaks (mandatory)
- Access to mental health support
- Option to skip distressing content
- Fair compensation for difficult work

Disagreement in Safety Annotation

Safety judgments are often subjective

Sources of disagreement:

- Cultural backgrounds
- Personal experiences
- Risk tolerance
- Context interpretation

Approaches:

- Use multiple annotators (3-5 per item)
- Report agreement alongside labels
- Consider disaggregated labels
- Weight by annotator demographics

Over-Refusal Problem

Safety vs. helpfulness tradeoff

Issue: Systems that refuse too often are unusable

Examples of over-refusal:

- Refusing medical questions
- Blocking historical discussions
- Rejecting creative writing prompts
- Refusing educational content about sensitive topics

Annotation task: Identify false positives

- Was this refusal appropriate?
- Should a helpful response have been given?

Same content, different intent

Challenge: Information can be used for good or harm

Examples:

- Security research vs. hacking
- Medical information vs. self-harm
- Chemistry education vs. dangerous synthesis

Annotation approach:

- Consider likely user intent
- Evaluate information availability elsewhere
- Assess potential for harm vs. benefit
- Flag for additional review when uncertain

Using LLMs to scale safety evaluation

Appropriate uses:

- Initial filtering of obvious cases
- Flagging content for human review
- Consistency checking

Limitations:

- May miss subtle harms
- Bias in training data
- Shouldn't be sole judge

Best practice: Human-in-the-loop for final decisions

Building Safety Datasets

Key principles:

- ① **Diverse harm coverage:** All relevant categories
- ② **Balanced data:** Not just obvious cases
- ③ **Edge cases:** Include borderline examples
- ④ **Multiple perspectives:** Diverse annotator pool
- ⑤ **Documentation:** Clear annotation guidelines
- ⑥ **Ethics review:** IRB or equivalent approval

Release considerations:

- Access restrictions for sensitive content
- Responsible use agreements

Next Class: Low-Resource Languages

Lecture 25 (Apr 27): Low-Resource Languages & Scaling

Topics:

- Challenges of low-resource annotation
- Cross-lingual transfer
- Native speaker recruitment
- Quality assurance at scale
- Case studies: MasakhaNER, AmericasNLU

Project: Final presentations May 4

Key Takeaways

- ① **Red teaming** finds vulnerabilities before deployment
- ② **Harmfulness taxonomies** structure safety annotation
- ③ **Severity scales** help prioritize issues
- ④ **Annotator wellbeing** is an ethical requirement
- ⑤ **Disagreement is expected** – use multiple annotators
- ⑥ **Balance safety with helpfulness** – avoid over-refusal

Questions?

Questions?

Office Hours: Wednesdays 1-3pm, Volen 109

 jinzhaob@brandeis.edu