

Best Practices & Future Directions

Course Summary and Looking Ahead

Jin Zhao

Brandeis University

April 29, 2025

Today's Agenda

- ① Course summary
- ② Best practices checklist
- ③ Emerging trends in annotation
- ④ The future of human annotation
- ⑤ Career paths
- ⑥ Final project presentations

Project: Final presentations May 4

Course Journey

What we've covered:

Foundations:

- Why annotation matters
- MATTER/MAMA cycles
- Task types
- Guidelines design

Practice:

- Annotation tools
- LLM annotation
- Human-AI collaboration

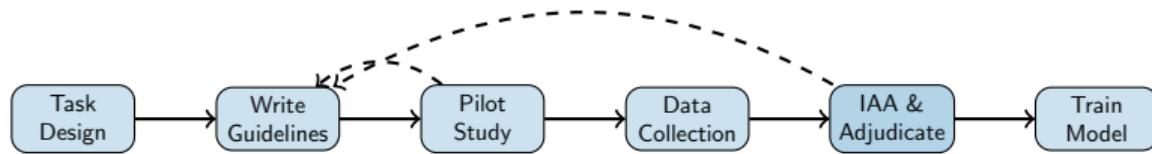
Quality:

- Inter-annotator agreement
- Adjudication
- Error analysis

Applications:

- Model training
- RLHF and preferences
- Safety annotation
- Low-resource settings

The Annotation Pipeline



Key insight: Annotation is iterative, not linear

Best Practices Checklist: Planning

Before you start:

- Define clear task objectives
- Choose appropriate task formalization
- Identify target annotator population
- Estimate data requirements
- Plan budget and timeline
- Select annotation tool
- Design evaluation metrics

Best Practices Checklist: Guidelines

Writing effective guidelines:

- Clear definitions for all categories
- Examples for each category
- Edge cases and how to handle them
- Decision trees for complex decisions
- What NOT to annotate
- Version control for updates
- Training materials

Best Practices Checklist: Quality

Ensuring annotation quality:

- Multiple annotators per item (2-3 minimum)
- Calculate and report IAA
- Regular calibration sessions
- Monitor per-annotator performance
- Embed gold standards for quality checks
- Document adjudication process
- Iterative guideline refinement

Best Practices Checklist: Ethics

Responsible annotation:

- Fair annotator compensation
- Clear terms and expectations
- Content warnings for sensitive material
- Mental health support if needed
- Data privacy protections
- Community consent for low-resource languages
- Proper attribution in publications

Best Practices Checklist: Documentation

For reproducibility:

- Data statement / datasheet
- Annotation guidelines (versioned)
- Annotator demographics
- IAA metrics and methodology
- Data format specification
- Known limitations
- License and usage terms

Common Mistakes to Avoid

Don't:

- ① Skip the pilot study
- ② Use only one annotator
- ③ Write vague guidelines
- ④ Ignore disagreements
- ⑤ Treat annotation as “just labeling”
- ⑥ Underpay annotators
- ⑦ Forget to document decisions
- ⑧ Train on test data

Emerging Trend: LLM-in-the-Loop

Human-AI collaboration is evolving

Current approaches:

- LLM pre-annotation with human correction
- LLM as “annotator 3” for tie-breaking
- Human review of LLM annotations

Emerging:

- Active learning with LLM uncertainty
- LLM-generated annotation guidelines
- Automated quality estimation
- LLM explanation of annotations

Emerging Trend: Synthetic Data

LLM-generated training data

Use cases:

- Data augmentation
- Rare category generation
- Privacy-preserving data

Challenges:

- Quality verification still needs humans
- Risk of bias amplification
- Not suitable for evaluation data

Future:

Hybrid human-synthetic datasets

Beyond classification

Growth areas:

- RLHF for model alignment
- DPO and alternatives
- Constitutional AI
- Multi-objective preferences

Annotation implications:

- New task types (comparison, ranking)
- Scalability challenges
- Subjectivity is a feature, not a bug
- Need for diverse annotator perspectives

The Future of Human Annotation

Will LLMs replace human annotators?

LLMs will take over:

- Simple, objective classification
- Large-scale pre-labeling
- Quality filtering

Humans remain essential for:

- Evaluation and benchmarking
- Subjective judgments
- Novel task design
- Safety-critical applications
- Low-resource languages
- Capturing diverse perspectives

Evolving Annotator Role

From labeler to expert reviewer

Traditional:

- Assign labels from scratch
- High volume, repetitive

Future:

- Review and correct AI predictions
- Handle edge cases and ambiguity
- Provide feedback on AI behavior
- Design and validate annotation schemes
- Train and calibrate AI systems

Higher skill, higher value

Career Paths in Annotation

Where this knowledge leads:

Industry roles:

- Data Operations Manager
- Annotation Quality Lead
- ML Data Specialist
- Trust & Safety Analyst
- Human-AI Interaction Designer

Research paths:

- Computational linguistics
- Human-computer interaction
- AI safety research
- Low-resource NLP

Companies Working on Annotation

The ecosystem:

AI companies (internal teams):

- OpenAI, Anthropic, Google, Meta, Microsoft
- Large data ops teams for RLHF

Annotation platforms:

- Scale AI, Labelbox, Appen, Surge AI

Tool providers:

- Label Studio, Prodigy, Argilla

Research labs and universities

May 4: Final Presentations

What to present:

- ① Task definition and motivation
- ② Annotation scheme and guidelines
- ③ Data collection process
- ④ IAA results and analysis
- ⑤ Model results (if applicable)
- ⑥ Lessons learned

Format: 15 minutes + 5 minutes Q&A

Final Report

Due: May 11

Contents:

- ① Executive summary**
- ② Task motivation and related work**
- ③ Annotation scheme design**
- ④ Data collection methodology**
- ⑤ Quality analysis (IAA, error analysis)**
- ⑥ Modeling results**
- ⑦ Discussion and future work**
- ⑧ Appendix: Guidelines, data samples**

Length: 8-12 pages (excluding appendix)

Key Takeaways from the Course

- ➊ **Annotation is fundamental** to ML – garbage in, garbage out
 - ➋ **Task design** determines annotation quality
 - ➌ **Good guidelines** are iteratively refined
 - ➍ **IAA measures** quality, not just agreement
 - ➎ **Human+AI** beats either alone
 - ➏ **Ethics matter** – annotators are people
 - ➐ **Documentation enables** reproducibility

Thank You!

Thank you for a great semester!

Good luck with your final presentations!

Office Hours: Available by appointment for project help

 jinzhaobr@brandeis.edu

Questions?

Questions?

Final Presentations: May 4

Final Report Due: May 11

 jinzhaob@brandeis.edu